

# Skin Cancer-main

April 25, 2024

```
[1]: #Import necessary libraries
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

import torchvision
from torchvision.datasets import ImageFolder
from torch.utils.data import Dataset, DataLoader
from torch.optim.lr_scheduler import ReduceLROnPlateau

from albumentations import (ToFloat, Normalize, VerticalFlip, HorizontalFlip,
    ↳Compose, Resize,
                                RandomBrightnessContrast, HueSaturationValue, Blur,
    ↳GaussNoise,
                                Rotate, RandomResizedCrop, ShiftScaleRotate)
from albumentations.pytorch import ToTensorV2
from torchsampler import ImbalancedDatasetSampler

from efficientnet_pytorch import EfficientNet

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import cv2

from sklearn.metrics import confusion_matrix, accuracy_score, precision_score,
    ↳recall_score, f1_score, roc_curve, auc, roc_auc_score, classification_report
from sklearn.preprocessing import LabelEncoder, normalize
from sklearn.model_selection import GroupKFold, train_test_split

from PIL import Image
import os
```

```
import itertools
import gc
import warnings
import random
from tqdm import tqdm

import datetime
from timeit import default_timer as timer
```

```
[2]: #Lets set up some things
device = 'cuda:0' if torch.cuda.is_available() else torch.device("cpu")
device
```

```
[2]: 'cuda:0'
```

```
[3]: warnings.simplefilter('ignore')
#Set up seeds
SEED = 22

#Set up seeds from reproducibility
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
```

## 0.1 Lets first Load the data and take a look at it

```
[4]: #Lets load the train and test data
df = pd.read_csv('./data/train.csv')
train_dir = './data/jpeg/train/'
```

```
[5]: df['target'].value_counts()
```

```
[5]: target
0    32542
1     584
Name: count, dtype: int64
```

```
[6]: #Since we don't have labels for the test set in our data
#Lets only use the train data and create training and testing set by splitting it

X = df.drop(columns=['target'])
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=SEED, stratify=y)
```

```

train_df = pd.concat([X_train, y_train], axis=1)
test_df = pd.concat([X_test, y_test], axis=1)
train_df.reset_index(drop=True, inplace=True)
test_df.reset_index(drop=True, inplace=True)

print("Train DataFrame shape:", train_df.shape)
print("Test DataFrame shape:", test_df.shape)

```

Train DataFrame shape: (26500, 8)  
Test DataFrame shape: (6626, 8)

```
[7]: train_df.head()
```

```

[7]:   image_name  patient_id  sex  age_approx  anatom_site_general_challenge \
0  ISIC_8104064  IP_7207550  male      55.0                        torso
1  ISIC_6917587  IP_0894335  female     40.0                      head/neck
2  ISIC_3391651  IP_2842809  female     55.0                        torso
3  ISIC_4547675  IP_7279968  male      45.0          upper extremity
4  ISIC_7289411  IP_5439716  male      70.0                        torso

   diagnosis  benign_malignant  target
0   unknown             benign        0
1   unknown             benign        0
2   unknown             benign        0
3   unknown             benign        0
4    nevus              benign        0

```

```
[8]: test_df.head()
```

```

[8]:   image_name  patient_id  sex  age_approx  anatom_site_general_challenge \
0  ISIC_0476762  IP_6323321  male      30.0          lower extremity
1  ISIC_1560888  IP_9147454  female     55.0          upper extremity
2  ISIC_1859923  IP_6420568  male      40.0                        torso
3  ISIC_9034411  IP_7517320  male      60.0                        torso
4  ISIC_5017874  IP_7517320  male      60.0          upper extremity

   diagnosis  benign_malignant  target
0    nevus              benign        0
1   unknown             benign        0
2   unknown             benign        0
3   unknown             benign        0
4   unknown             benign        0

```

We can see that there are around 33k images for train data and 10k images for test data.

```
[9]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 26500 entries, 0 to 26499

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	image_name	26500 non-null	object
1	patient_id	26500 non-null	object
2	sex	26449 non-null	object
3	age_approx	26446 non-null	float64
4	anatom_site_general_challenge	26084 non-null	object
5	diagnosis	26500 non-null	object
6	benign_malignant	26500 non-null	object
7	target	26500 non-null	int64

dtypes: float64(1), int64(1), object(6)

memory usage: 1.6+ MB

```
[10]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 6626 entries, 0 to 6625

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	image_name	6626 non-null	object
1	patient_id	6626 non-null	object
2	sex	6612 non-null	object
3	age_approx	6612 non-null	float64
4	anatom_site_general_challenge	6515 non-null	object
5	diagnosis	6626 non-null	object
6	benign_malignant	6626 non-null	object
7	target	6626 non-null	int64

dtypes: float64(1), int64(1), object(6)

memory usage: 414.2+ KB


```
[11]: def plot_missing(df, data):  
    missing_counts = df.isna().sum()  
    missing_counts_df = pd.DataFrame(missing_counts, columns=['Null Values'])  
    plt.figure(figsize=(10,6))  
    sns.heatmap(missing_counts_df, cmap="YlGnBu", annot=True, fmt='g')  
    plt.title(f"Count of Null Values in {data} data")  
    plt.xlabel('Columns')  
    plt.ylabel('Counts')  
    plt.xticks(rotation=90)  
    plt.show()
```

```
[12]: train_df.isna().sum()
```

```
[12]: image_name          0
      patient_id         0
      sex                51
      age_approx         54
      anatom_site_general_challenge  416
      diagnosis          0
      benign_malignant    0
      target             0
      dtype: int64
```

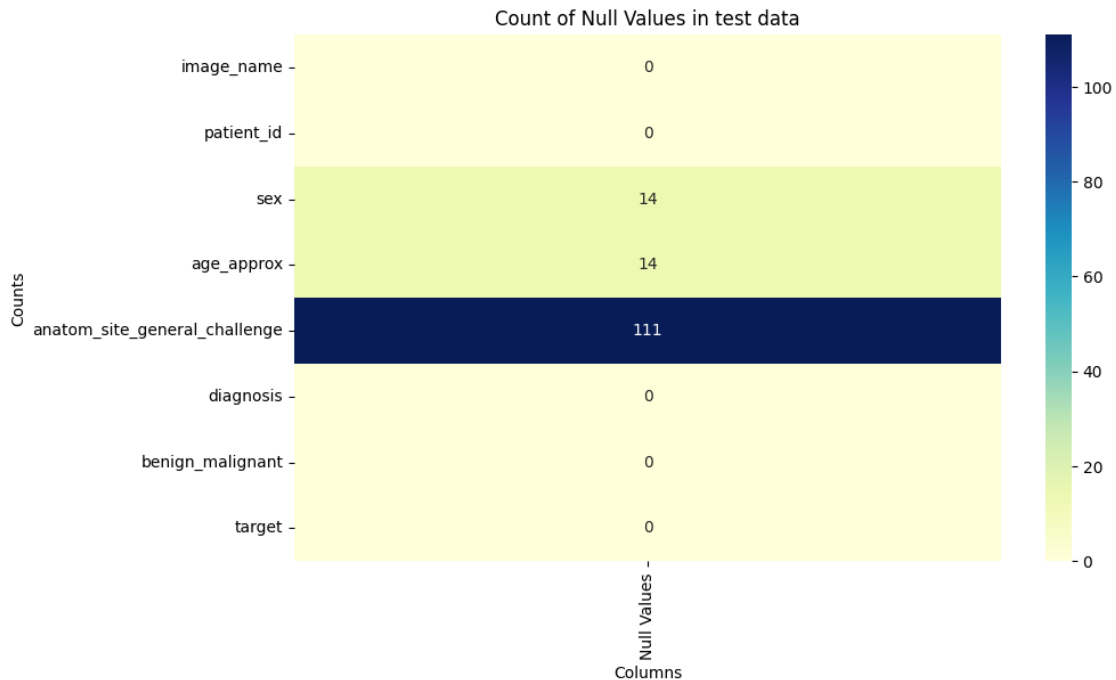
## 0.2 Missing Values for Train Data

```
[13]: plot_missing(train_df, "train")
```



## 0.3 Missing Values for Test Data

```
[14]: plot_missing(test_df, "test")
```



```
[15]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6626 entries, 0 to 6625
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   image_name                            6626 non-null   object
1   patient_id                            6626 non-null   object
2   sex                                    6612 non-null   object
3   age_approx                            6612 non-null   float64
4   anatom_site_general_challenge         6515 non-null   object
5   diagnosis                             6626 non-null   object
6   benign_malignant                      6626 non-null   object
7   target                                6626 non-null   int64
dtypes: float64(1), int64(1), object(6)
memory usage: 414.2+ KB
```

```
[16]: null_counts_train = train_df.groupby('target').apply(lambda x: x.isnull().sum())
null_counts_train
```

```
[16]:      image_name  patient_id  sex  age_approx  \
target
0           0           0    51          54
1           0           0     0           0
```

	anatom_site_general_challenge	diagnosis	benign_malignant	target
target				
0	408	0	0	0
1	8	0	0	0

```
[17]: null_counts_test = test_df.groupby('target').apply(lambda x: x.isnull().sum())
null_counts_test
```


```
[17]: image_name  patient_id  sex  age_approx  \
target
0          0          0  14          14
1          0          0   0           0
```

	anatom_site_general_challenge	diagnosis	benign_malignant	target
target				
0	110	0	0	0
1	1	0	0	0


From above we can see that in training data *sex*, *age\_approx* and *anatom\_site\_general\_challenge* have missing values and in test data *anatom\_site\_general\_challenge* has missing value. Since most of them belong to benign class lets drop them.

```
[18]: train_df = train_df.dropna()
test_df = test_df.dropna()
```

```
[19]: plot_missing(train_df, "train")
```



```
[20]: plot_missing(test_df, "test")
```




Lets Plot the Malignent vs Benign case

```
[21]: malignant = len(train_df[train_df["target"] == 1])
benign = len(train_df[train_df["target"] == 0])
colors = ['#ff9999', '#99ff99']
labels = ["Malignant", "Benign"]
size = [malignant, benign]

plt.figure(figsize = (5, 5))
plt.pie(size, labels = labels, explode=[0.1,0.0], shadow = True, startangle = 90, colors = colors, autopct='%1.1f%%')
plt.title("Malignant VS Benign Cases")
plt.legend()
```

```
[21]: <matplotlib.legend.Legend at 0x75070a0e3100>
```





The data is imbalanced.

```
[22]: males_count_tr = len(train_df[train_df['sex']=='male'])
      females_count_tr = len(train_df[train_df['sex']=="female"])

      males_count_ts = len(test_df[test_df['sex']=='male'])
      females_count_ts = len(test_df[test_df['sex']=="female"])


      labels = ['Males', 'Females']
      size = [males_count_tr, females_count_tr]
      explode = [0.1, 0.0]

      plt.figure(figsize= (10,10))
      plt.subplot(1,2,1)
      plt.pie(size, labels=labels, explode=explode, shadow=True, startangle=90,
              colors= colors, autopct='%1.1f%%')
      plt.title('Male vs Female Training Set Count')
      plt.legend()

      size = [males_count_ts, females_count_ts]
      plt.subplot(1,2,2)
```

```
plt.pie(size, labels = labels, explode = explode, shadow = True, startangle = 90, colors = colors, autopct='%1.1f%%')
plt.title("Male vs Female Test Set Count")
plt.legend()
```

[22]: <matplotlib.legend.Legend at 0x75070a14d100>

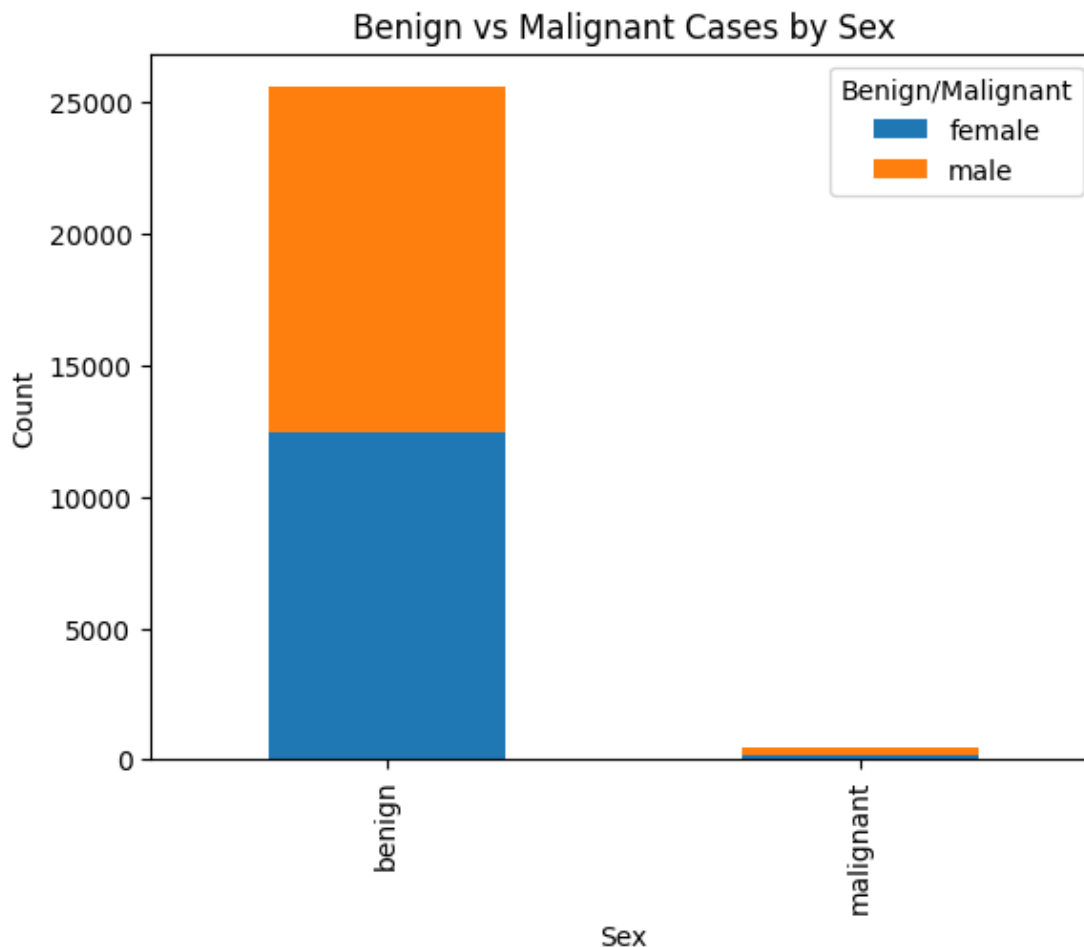


#### 0.4 Benign and malignant by sex.

```
[23]: x = train_df.groupby(['benign_malignant', 'sex']).size()
x
```

```
[23]: benign_malignant  sex
benign                female    12413
                   male        13158
malignant             female      175
                   male         284
dtype: int64
```

```
[24]: x.unstack().plot(kind='bar', stacked=True)
plt.xlabel('Sex')
plt.ylabel('Count')
plt.title('Benign vs Malignant Cases by Sex')
plt.legend(title='Benign/Malignant')
plt.show()
```



```
[25]: print("##### Training set info #####")
print(train_df['anatom_site_general_challenge'].unique())
print(train_df['anatom_site_general_challenge'].value_counts())

print("\n\n")

print("##### Test set info #####")
print(test_df['anatom_site_general_challenge'].unique())
print(test_df['anatom_site_general_challenge'].value_counts())
```

```
##### Training set info #####
['torso' 'head/neck' 'upper extremity' 'lower extremity' 'palms/soles'
 'oral/genital']
anatom_site_general_challenge
torso          13423
lower extremity    6754
upper extremity    3970
```

```

head/neck          1477
palms/soles        298
oral/genital       108
Name: count, dtype: int64

```

```

##### Test set info #####
['lower extremity' 'upper extremity' 'torso' 'palms/soles' 'head/neck'
 'oral/genital']
anatom_site_general_challenge
torso          3402
lower extremity 1645
upper extremity  993
head/neck       368
palms/soles      77
oral/genital     16
Name: count, dtype: int64

```

Let's plot the distribution for each columns in train data.

We can see that majority of the cases are observed at the torso, and after that the upper and lower extremities of the body in both the training and testing set.

```


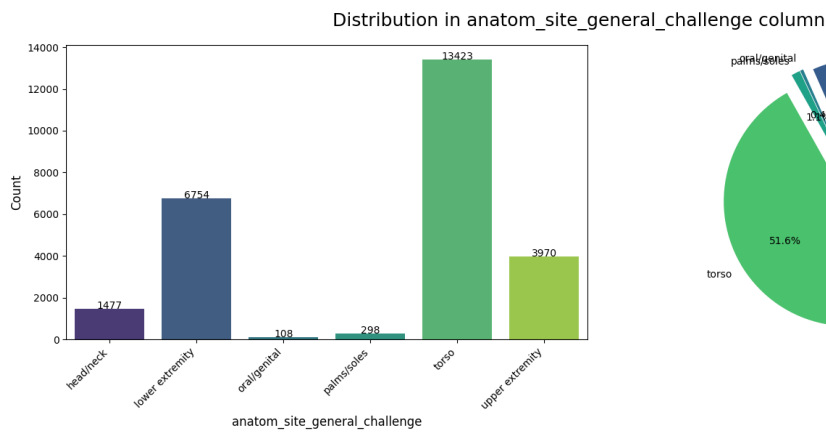




[26]: def plot_columns(df, column):
        value_counts = df[column].value_counts().sort_index()
        fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))
        sns.barplot(x=value_counts.index, y=value_counts.values, palette="viridis",
        ↪ax=ax1)
        ax1.set_xlabel(column, fontsize=12)
        ax1.set_ylabel('Count', fontsize=12)
        ax1.set_xticklabels(ax1.get_xticklabels(), rotation=45, ha='right',
        ↪fontsize=10)
        for index, value in enumerate(value_counts.values):
            ax1.text(index, value + 0.1, str(value), ha='center', fontsize=10)
        explode = [0.1 for i in range(len(value_counts))]
        ax2.pie(value_counts, labels=value_counts.index, autopct='%1.1f%%',
        ↪colors=sns.color_palette("viridis", len(value_counts)), explode=explode)
        ax2.axis('equal')
        fig.suptitle(f'Distribution in {column} column', fontsize=18)
        plt.tight_layout()
        plt.show()

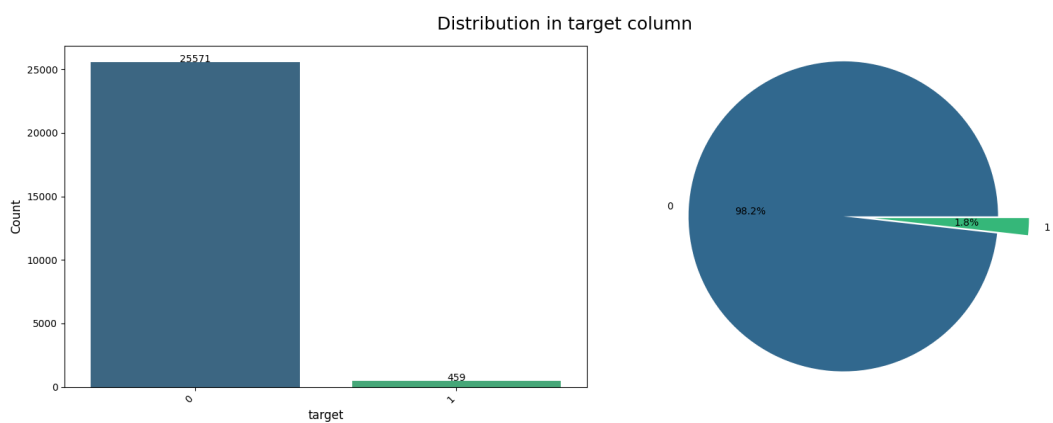

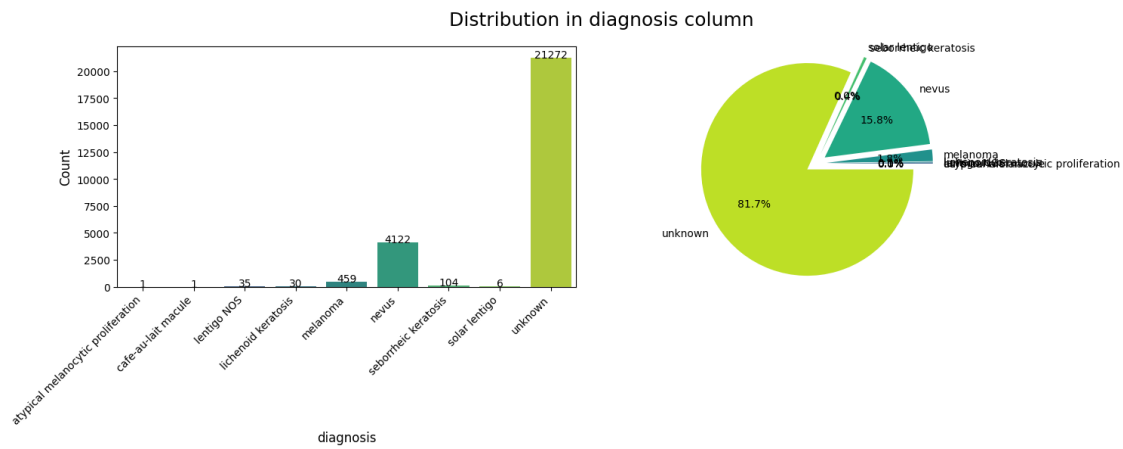
```

```

[27]: for column in train_df.columns:
        if column == 'image_name' or column == 'patient_id':
            continue
        plot_columns(train_df, column)




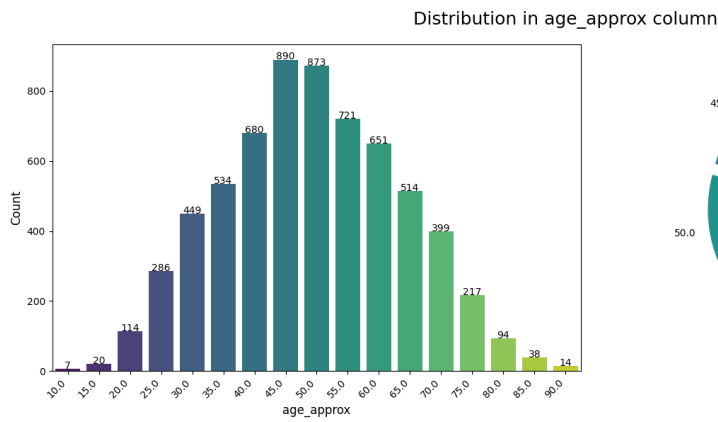


```

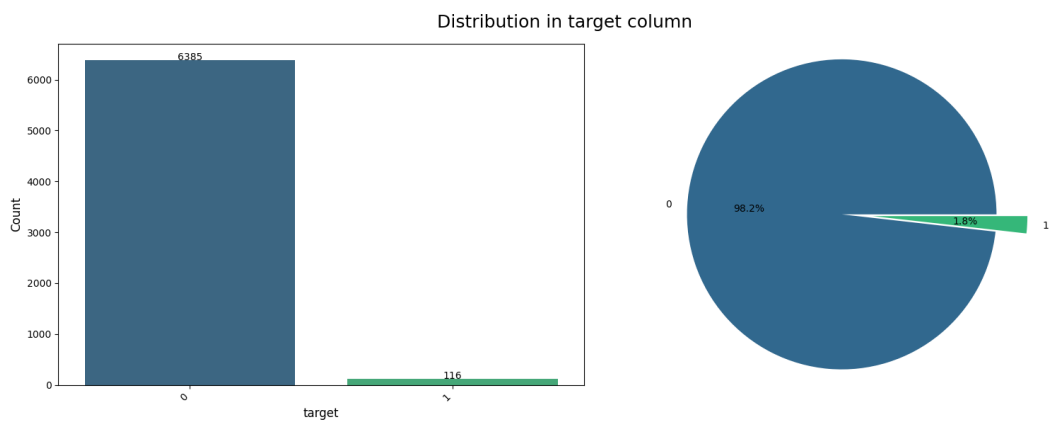
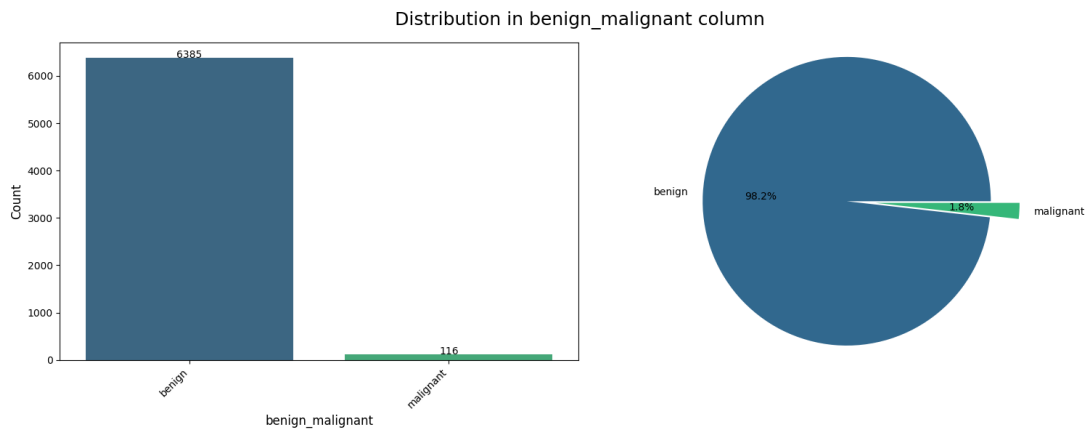





```
[28]: for column in test_df.columns:
        if column == 'image_name' or column == 'patient_id':
```

```
continue
plot_columns(test_df, column)
```






Let's view some images from our data.



```
[29]: sample = train_df.sample(12).reset_index()
images = [sample['image_name'][i] + '.jpg' for i in range(12)]
labels = [sample['target'][i] for i in range(12)]
d_map = {1:"malignant",0:"benign"}
plt.figure(figsize = (16,6))
for i in range(12):
    plt.subplot(2,6, i+1)
    img = cv2.imread(os.path.join(train_dir, images[i]))
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    plt.title(d_map[labels[i]])
    plt.axis('off')
    plt.imshow(img, cmap='gray')
```



## 0.5 Lets Standardize and preprocess the data.

```
[30]: def drop_and_encode(df, type):
    to_drop = ['diagnosis', 'benign_malignant']
    for col in to_drop:
        if col in df.columns:
            df.drop([col], axis=1,inplace=True)

    to_encode = ['sex', 'anatom_site_general_challenge']
    encoded_all = []

    #Lets encode sex and anatom_site_general_challenge columns
    df[to_encode[0]] = df[to_encode[0]].astype(str)
    df[to_encode[1]] = df[to_encode[1]].astype(str)

    label_encoder = LabelEncoder()

    for col in to_encode:
```

```

        encoded = label_encoder.fit_transform(df[col])
        encoded_all.append(encoded)

    label_mappings = dict(zip(label_encoder.classes_, label_encoder.
↪transform(label_encoder.classes_)))
    print(f"Label_mappings for anatom_site_general_challenge:\n",label_mappings)

    df[to_encode[0]] = encoded_all[0]
    df[to_encode[1]] = encoded_all[1]

    #Add image path to the dataframe
    if type == "train":
        df['image_name'] = '/home/pukar/Pictures/skin-cancer/data/jpeg/train/' +
↪train_df['image_name'] + '.jpg'
    if type == "test":
        df['image_name'] = '/home/pukar/Pictures/skin-cancer/data/jpeg/train/' +
↪test_df['image_name'] + '.jpg'

```

```

[31]: drop_and_encode(train_df,type="train")
      drop_and_encode(test_df,type="test")

```

```

Label_mappings for anatom_site_general_challenge:
{'head/neck': 0, 'lower extremity': 1, 'oral/genital': 2, 'palms/soles': 3,
'torso': 4, 'upper extremity': 5}
Label_mappings for anatom_site_general_challenge:
{'head/neck': 0, 'lower extremity': 1, 'oral/genital': 2, 'palms/soles': 3,
'torso': 4, 'upper extremity': 5}

```

```

[32]: #Lets remane some columns
train_df = train_df.rename(columns={'age_approx': 'age',
↪'anatom_site_general_challenge': 'anatomy'})
test_df = test_df.rename(columns={'age_approx': 'age',
↪'anatom_site_general_challenge': 'anatomy'})

train_df.reset_index(drop=True, inplace=True)
test_df.reset_index(drop=True, inplace=True)
train_df.head()

```

```

[32]:
           image_name  patient_id  sex  age  \
0  /home/pukar/Pictures/skin-cancer/data/jpeg/tra...  IP_7207550    1  55.0
1  /home/pukar/Pictures/skin-cancer/data/jpeg/tra...  IP_0894335    0  40.0
2  /home/pukar/Pictures/skin-cancer/data/jpeg/tra...  IP_2842809    0  55.0
3  /home/pukar/Pictures/skin-cancer/data/jpeg/tra...  IP_7279968    1  45.0
4  /home/pukar/Pictures/skin-cancer/data/jpeg/tra...  IP_5439716    1  70.0

      anatomy  target
0          4        0

```

1	0	0
2	4	0
3	5	0
4	4	0

```
[33]: #Normalize the columns
normalized_train = normalize(train_df[['sex', 'age', 'anatomy']])
normalized_test = normalize(test_df[['sex', 'age', 'anatomy']])

train_df['sex'] = normalized_train[:, 0]
train_df['age'] = normalized_train[:, 1]
train_df['anatomy'] = normalized_train[:, 2]

test_df['sex'] = normalized_test[:, 0]
test_df['age'] = normalized_test[:, 1]
test_df['anatomy'] = normalized_test[:, 2]

print('Len Train: {:,}'.format(len(train_df)), '\n' +
      'Len Test: {:,}'.format(len(test_df)))
```

Len Train: 26,030

Len Test: 6,501

```
[34]: train_df['target'].value_counts()
```

```
[34]: target
0    25571
1      459
Name: count, dtype: int64
```

```
[35]: test_df['target'].value_counts()
```

```
[35]: target
0    6385
1     116
Name: count, dtype: int64
```

## 0.6 Now Let's Create Custom Dataset Class and Data Loaders

```
[36]: feat_cols = ['sex', 'age', 'anatomy']
no_cols = 3
output_size = 1
```

```
[37]: class MelanomaDataset(Dataset):
      def __init__(self, dataframe, is_train=True, is_val=False, is_test=False):
```

```

        self.dataframe, self.is_train, self.is_valid, self.is_test = dataframe,
↪is_train, is_val, is_test
        self.horizontal_flip = 0.5
        self.vertical_flip = 0.5
        if is_train or is_test:
            self.transform = Compose([RandomResizedCrop(height=224, width=224,
↪scale=(0.4, 1.0)),
                                ShiftScaleRotate(rotate_limit=90,
↪scale_limit = [0.8, 1.2]),
                                HorizontalFlip(p = self.horizontal_flip),
                                VerticalFlip(p = self.vertical_flip),
                                HueSaturationValue(sat_shift_limit=[0.7,
↪1.3],
                                hue_shift_limit=[-0.1,
↪0.1]),
                                RandomBrightnessContrast(brightness_limit=[0.7, 1.3],
                                contrast_limit=
↪[0.7, 1.3]),
                                Normalize(),
                                ToTensorV2())])
        else:
            self.transform = Compose([
                                Resize(224,224),
                                Normalize(),
                                ToTensorV2())])

    def __len__(self):
        return len(self.dataframe)

    def __getitem__(self, idx):
        img_pth = self.dataframe['image_name'][idx]
        image = cv2.imread(img_pth)
        col_features = np.array(self.dataframe.iloc[idx][['sex', 'age',
↪'anatomy']]).values, dtype=np.float32)

        #Apply transforms
        image = self.transform(image = image)
        image = image['image']

        return image, col_features, self.dataframe['target'][idx]

    def get_labels(self): return self.dataframe['target']

```

```
[38]: def dataset_collate(batch):
    images = []
    col_features = []
    targets = []

    for image, col_feature, target in batch:
        images.append(image)
        col_features.append(col_feature)
        targets.append(target)
    print(images[0].shape)
    images = np.array(images)
    col_features = np.array(col_features)
    targets = np.array(targets)
    return images, col_features, targets
```

```
[39]: # Test Data and Loader and sampler
example_data = MelanomaDataset(train_df, is_train=False, is_val=True,
    ↪is_test=False)
example_loader = torch.utils.data.DataLoader(example_data, batch_size = 5,
    ↪sampler=ImbalancedDatasetSampler(example_data))

# Get a sample
for k, (images, csv_data, lab) in enumerate(example_loader):
    images = torch.tensor(images, device=device, dtype=torch.float32)
    csv_data = torch.tensor(csv_data, device=device, dtype=torch.float32)
    print(k, csv_data)
    print(images.shape)
    print(lab)
    break
    #valid_preds[k*images.shape[0] : k*images.shape[0] + images.shape[0]]
```

```
0 tensor([[0.0181, 0.9972, 0.0725],
          [0.0154, 0.9980, 0.0614],
          [0.0249, 0.9947, 0.0995],
          [0.0000, 1.0000, 0.0000],
          [0.0133, 0.9985, 0.0533]], device='cuda:0')
torch.Size([5, 3, 224, 224])
tensor([1, 1, 0, 1, 1])
```

```
[40]: class MelanomaClassifier(nn.Module):
    def __init__(self, output_size, no_cols: int):
        super(MelanomaClassifier, self).__init__()

        model_name = f"efficientnet-b2"
```

```

#For image features
self.cnn = EfficientNet.from_pretrained(model_name)

#For col features
self.cols = nn.Sequential(
    nn.Linear(no_cols, 250),
    nn.BatchNorm1d(250),
    nn.ReLU(),
    nn.Dropout(p=0.3),

    nn.Linear(250, 250),
    nn.BatchNorm1d(250),
    nn.ReLU(),
    nn.Dropout(p=0.3)
)

#EfficientNet B2 outputs 1408 freatures, now lets change its FCN for
↪ classification
self.classifier = nn.Sequential(nn.Linear(1408 + 250, output_size))

def forward(self, image, col_features):

    #Feed image to CNN to extract features
    image = self.cnn.extract_features(image)
    image = F.avg_pool2d(image, image.size()[2:]).reshape(-1, 1408)

    #Feed col features to the FNN
    col_features = self.cols(col_features)

    #Combine two features
    combined_features = torch.cat((image, col_features), dim=1)

    #Use classifier model
    output = self.classifier(combined_features)

    return output

```

```

[41]: M= MelanomaClassifier(output_size=output_size, no_cols=no_cols)
      M

```

Loaded pretrained weights for efficientnet-b2

```

[41]: MelanomaClassifier(
      (cnn): EfficientNet(
        (_conv_stem): Conv2dStaticSamePadding(
          3, 32, kernel_size=(3, 3), stride=(2, 2), bias=False
        (static_padding): ZeroPad2d((0, 1, 0, 1))

```

```

)
(_bn0): BatchNorm2d(32, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
(_blocks): ModuleList(
  (0): MBConvBlock(
    (_depthwise_conv): Conv2dStaticSamePadding(
      32, 32, kernel_size=(3, 3), stride=[1, 1], groups=32, bias=False
      (static_padding): ZeroPad2d((1, 1, 1, 1))
    )
    (_bn1): BatchNorm2d(32, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_se_reduce): Conv2dStaticSamePadding(
      32, 8, kernel_size=(1, 1), stride=(1, 1)
      (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
      8, 32, kernel_size=(1, 1), stride=(1, 1)
      (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
      32, 16, kernel_size=(1, 1), stride=(1, 1), bias=False
      (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(16, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
  )
  (1): MBConvBlock(
    (_depthwise_conv): Conv2dStaticSamePadding(
      16, 16, kernel_size=(3, 3), stride=(1, 1), groups=16, bias=False
      (static_padding): ZeroPad2d((1, 1, 1, 1))
    )
    (_bn1): BatchNorm2d(16, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_se_reduce): Conv2dStaticSamePadding(
      16, 4, kernel_size=(1, 1), stride=(1, 1)
      (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
      4, 16, kernel_size=(1, 1), stride=(1, 1)
      (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
      16, 16, kernel_size=(1, 1), stride=(1, 1), bias=False
      (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(16, eps=0.001, momentum=0.010000000000000009,

```

```

affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
)
(2): MBConvBlock(
  (_expand_conv): Conv2dStaticSamePadding(
    16, 96, kernel_size=(1, 1), stride=(1, 1), bias=False
  (static_padding): Identity()
  )
  (_bn0): BatchNorm2d(96, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_depthwise_conv): Conv2dStaticSamePadding(
    96, 96, kernel_size=(3, 3), stride=[2, 2], groups=96, bias=False
  (static_padding): ZeroPad2d((0, 1, 0, 1))
  )
  (_bn1): BatchNorm2d(96, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_se_reduce): Conv2dStaticSamePadding(
    96, 4, kernel_size=(1, 1), stride=(1, 1)
  (static_padding): Identity()
  )
  (_se_expand): Conv2dStaticSamePadding(
    4, 96, kernel_size=(1, 1), stride=(1, 1)
  (static_padding): Identity()
  )
  (_project_conv): Conv2dStaticSamePadding(
    96, 24, kernel_size=(1, 1), stride=(1, 1), bias=False
  (static_padding): Identity()
  )
  (_bn2): BatchNorm2d(24, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_swish): MemoryEfficientSwish()
)
(3-4): 2 x MBConvBlock(
  (_expand_conv): Conv2dStaticSamePadding(
    24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False
  (static_padding): Identity()
  )
  (_bn0): BatchNorm2d(144, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_depthwise_conv): Conv2dStaticSamePadding(
    144, 144, kernel_size=(3, 3), stride=(1, 1), groups=144, bias=False
  (static_padding): ZeroPad2d((1, 1, 1, 1))
  )
  (_bn1): BatchNorm2d(144, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_se_reduce): Conv2dStaticSamePadding(
    144, 6, kernel_size=(1, 1), stride=(1, 1)

```



```

        (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
        6, 144, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
        144, 24, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(24, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
)
(5): MBConvBlock(
    (_expand_conv): Conv2dStaticSamePadding(
        24, 144, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn0): BatchNorm2d(144, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_depthwise_conv): Conv2dStaticSamePadding(
        144, 144, kernel_size=(5, 5), stride=[2, 2], groups=144, bias=False
        (static_padding): ZeroPad2d((2, 2, 2, 2))
    )
    (_bn1): BatchNorm2d(144, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_se_reduce): Conv2dStaticSamePadding(
        144, 6, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
        6, 144, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
        144, 48, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(48, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
)
(6-7): 2 x MBConvBlock(
    (_expand_conv): Conv2dStaticSamePadding(
        48, 288, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()

```

```

    )
    (_bn0): BatchNorm2d(288, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_depthwise_conv): Conv2dStaticSamePadding(
    288, 288, kernel_size=(5, 5), stride=(1, 1), groups=288, bias=False
    (static_padding): ZeroPad2d((2, 2, 2, 2))
    )
    (_bn1): BatchNorm2d(288, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_se_reduce): Conv2dStaticSamePadding(
    288, 12, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
    12, 288, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
    288, 48, kernel_size=(1, 1), stride=(1, 1), bias=False
    (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(48, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
    )
    (8): MBConvBlock(
    (_expand_conv): Conv2dStaticSamePadding(
    48, 288, kernel_size=(1, 1), stride=(1, 1), bias=False
    (static_padding): Identity()
    )
    (_bn0): BatchNorm2d(288, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_depthwise_conv): Conv2dStaticSamePadding(
    288, 288, kernel_size=(3, 3), stride=[2, 2], groups=288, bias=False
    (static_padding): ZeroPad2d((1, 1, 1, 1))
    )
    (_bn1): BatchNorm2d(288, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_se_reduce): Conv2dStaticSamePadding(
    288, 12, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
    12, 288, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(

```

```

        288, 88, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(88, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
)
(9-11): 3 x MBConvBlock(
    (_expand_conv): Conv2dStaticSamePadding(
        88, 528, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn0): BatchNorm2d(528, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_depthwise_conv): Conv2dStaticSamePadding(
        528, 528, kernel_size=(3, 3), stride=(1, 1), groups=528, bias=False
        (static_padding): ZeroPad2d((1, 1, 1, 1))
    )
    (_bn1): BatchNorm2d(528, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_se_reduce): Conv2dStaticSamePadding(
        528, 22, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
    )
    (_se_expand): Conv2dStaticSamePadding(
        22, 528, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
        528, 88, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(88, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
)
(12): MBConvBlock(
    (_expand_conv): Conv2dStaticSamePadding(
        88, 528, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn0): BatchNorm2d(528, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_depthwise_conv): Conv2dStaticSamePadding(
        528, 528, kernel_size=(5, 5), stride=[1, 1], groups=528, bias=False
        (static_padding): ZeroPad2d((2, 2, 2, 2))
    )
)

```

```

        (_bn1): BatchNorm2d(528, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_se_reduce): Conv2dStaticSamePadding(
          528, 22, kernel_size=(1, 1), stride=(1, 1)
          (static_padding): Identity()
        )
        (_se_expand): Conv2dStaticSamePadding(
          22, 528, kernel_size=(1, 1), stride=(1, 1)
          (static_padding): Identity()
        )
        (_project_conv): Conv2dStaticSamePadding(
          528, 120, kernel_size=(1, 1), stride=(1, 1), bias=False
          (static_padding): Identity()
        )
        (_bn2): BatchNorm2d(120, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_swish): MemoryEfficientSwish()
      )
    (13-15): 3 x MBConvBlock(
      (_expand_conv): Conv2dStaticSamePadding(
        120, 720, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
      )
      (_bn0): BatchNorm2d(720, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
      (_depthwise_conv): Conv2dStaticSamePadding(
        720, 720, kernel_size=(5, 5), stride=(1, 1), groups=720, bias=False
        (static_padding): ZeroPad2d((2, 2, 2, 2))
      )
      (_bn1): BatchNorm2d(720, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
      (_se_reduce): Conv2dStaticSamePadding(
        720, 30, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
      )
      (_se_expand): Conv2dStaticSamePadding(
        30, 720, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
      )
      (_project_conv): Conv2dStaticSamePadding(
        720, 120, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
      )
      (_bn2): BatchNorm2d(120, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
      (_swish): MemoryEfficientSwish()
    )

```

```

(16): MBConvBlock(
  (_expand_conv): Conv2dStaticSamePadding(
    120, 720, kernel_size=(1, 1), stride=(1, 1), bias=False
    (static_padding): Identity()
  )
  (_bn0): BatchNorm2d(720, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_depthwise_conv): Conv2dStaticSamePadding(
    720, 720, kernel_size=(5, 5), stride=[2, 2], groups=720, bias=False
    (static_padding): ZeroPad2d((2, 2, 2, 2))
  )
  (_bn1): BatchNorm2d(720, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_se_reduce): Conv2dStaticSamePadding(
    720, 30, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
  )
  (_se_expand): Conv2dStaticSamePadding(
    30, 720, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
  )
  (_project_conv): Conv2dStaticSamePadding(
    720, 208, kernel_size=(1, 1), stride=(1, 1), bias=False
    (static_padding): Identity()
  )
  (_bn2): BatchNorm2d(208, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_swish): MemoryEfficientSwish()
)
(17-20): 4 x MBConvBlock(
  (_expand_conv): Conv2dStaticSamePadding(
    208, 1248, kernel_size=(1, 1), stride=(1, 1), bias=False
    (static_padding): Identity()
  )
  (_bn0): BatchNorm2d(1248, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_depthwise_conv): Conv2dStaticSamePadding(
    1248, 1248, kernel_size=(5, 5), stride=(1, 1), groups=1248, bias=False
    (static_padding): ZeroPad2d((2, 2, 2, 2))
  )
  (_bn1): BatchNorm2d(1248, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
  (_se_reduce): Conv2dStaticSamePadding(
    1248, 52, kernel_size=(1, 1), stride=(1, 1)
    (static_padding): Identity()
  )
  (_se_expand): Conv2dStaticSamePadding(

```

```

        52, 1248, kernel_size=(1, 1), stride=(1, 1)
        (static_padding): Identity()
    )
    (_project_conv): Conv2dStaticSamePadding(
        1248, 208, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn2): BatchNorm2d(208, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_swish): MemoryEfficientSwish()
    )
    (21): MBConvBlock(
        (_expand_conv): Conv2dStaticSamePadding(
            208, 1248, kernel_size=(1, 1), stride=(1, 1), bias=False
            (static_padding): Identity()
        )
        (_bn0): BatchNorm2d(1248, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_depthwise_conv): Conv2dStaticSamePadding(
            1248, 1248, kernel_size=(3, 3), stride=[1, 1], groups=1248, bias=False
            (static_padding): ZeroPad2d((1, 1, 1, 1))
        )
        (_bn1): BatchNorm2d(1248, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_se_reduce): Conv2dStaticSamePadding(
            1248, 52, kernel_size=(1, 1), stride=(1, 1)
            (static_padding): Identity()
        )
        (_se_expand): Conv2dStaticSamePadding(
            52, 1248, kernel_size=(1, 1), stride=(1, 1)
            (static_padding): Identity()
        )
        (_project_conv): Conv2dStaticSamePadding(
            1248, 352, kernel_size=(1, 1), stride=(1, 1), bias=False
            (static_padding): Identity()
        )
        (_bn2): BatchNorm2d(352, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_swish): MemoryEfficientSwish()
    )
    (22): MBConvBlock(
        (_expand_conv): Conv2dStaticSamePadding(
            352, 2112, kernel_size=(1, 1), stride=(1, 1), bias=False
            (static_padding): Identity()
        )
        (_bn0): BatchNorm2d(2112, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)

```

```

        (_depthwise_conv): Conv2dStaticSamePadding(
            2112, 2112, kernel_size=(3, 3), stride=(1, 1), groups=2112, bias=False
            (static_padding): ZeroPad2d((1, 1, 1, 1))
        )
        (_bn1): BatchNorm2d(2112, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_se_reduce): Conv2dStaticSamePadding(
            2112, 88, kernel_size=(1, 1), stride=(1, 1)
            (static_padding): Identity()
        )
        (_se_expand): Conv2dStaticSamePadding(
            88, 2112, kernel_size=(1, 1), stride=(1, 1)
            (static_padding): Identity()
        )
        (_project_conv): Conv2dStaticSamePadding(
            2112, 352, kernel_size=(1, 1), stride=(1, 1), bias=False
            (static_padding): Identity()
        )
        (_bn2): BatchNorm2d(352, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
        (_swish): MemoryEfficientSwish()
    )
)
    (_conv_head): Conv2dStaticSamePadding(
        352, 1408, kernel_size=(1, 1), stride=(1, 1), bias=False
        (static_padding): Identity()
    )
    (_bn1): BatchNorm2d(1408, eps=0.001, momentum=0.010000000000000009,
affine=True, track_running_stats=True)
    (_avg_pooling): AdaptiveAvgPool2d(output_size=1)
    (_dropout): Dropout(p=0.3, inplace=False)
    (_fc): Linear(in_features=1408, out_features=1000, bias=True)
    (_swish): MemoryEfficientSwish()
)
(cols): Sequential(
  (0): Linear(in_features=3, out_features=250, bias=True)
  (1): BatchNorm1d(250, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): ReLU()
  (3): Dropout(p=0.3, inplace=False)
  (4): Linear(in_features=250, out_features=250, bias=True)
  (5): BatchNorm1d(250, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (6): ReLU()
  (7): Dropout(p=0.3, inplace=False)
)
(classifier): Sequential(

```

```

        (0): Linear(in_features=1658, out_features=1, bias=True)
    )
)

```

```

[42]: train_len = len(train_df)
      test_len = len(test_df)

      epochs = 30
      patience = 4
      TTA = 5
      k = 6

      weight_decay = 0.0
      num_workers = 2
      learning_rate = 0.0001
      lr_patience = 1
      lr_factor = 0.3

      batch_size = 32

      oof = np.zeros(shape = (train_len, 1))
      print('Out of Fold shape:', oof.shape, '\n')

      group_fold = GroupKFold(n_splits = k)

      # Generate index to split data into training and test set.
      folds = group_fold.split(X = np.zeros(train_len), y = train_df['target'],
                               groups = train_df['patient_id'].tolist())

```

Out of Fold shape: (26030, 1)

```

[43]: def train_folds(model):
      # Log file
      f = open(f"logs.txt", "w+")

      results = []
      for fold, (train_index, valid_index) in enumerate(folds):
          with open(f"logs.txt", 'a+') as f:
              print('-'*10, 'Fold:', fold+1, '-'*10, file=f)
              print('-'*10, 'Fold:', fold+1, '-'*10)

          # Best ROC score in this fold
          best_roc = None
          patience_f = patience

```



```

    # Reinitialize the model for every fold
    model = model

    optimizer = torch.optim.Adam(model.parameters(), lr = learning_rate,
↪weight_decay=weight_decay)
    scheduler = ReduceLROnPlateau(optimizer=optimizer, mode='max',
                                patience=lr_patience, verbose=True,
↪factor=lr_factor)
    criterion = nn.BCEWithLogitsLoss()

    # Train and Valid Data
    train_data = train_df.iloc[train_index].reset_index(drop=True)
    valid_data = train_df.iloc[valid_index].reset_index(drop=True)

    # Create Data instances
    train = MelanomaDataset(train_data, is_train=True, is_val=False,
↪is_test=False)
    valid = MelanomaDataset(valid_data, is_train=False, is_val=True,
↪is_test=False)
    test = MelanomaDataset(test_df, is_train=False, is_val=False,
↪is_test=True)

    # Create dataloader
    train_loader = DataLoader(train, batch_size=32,
↪sampler=ImbalancedDatasetSampler(train), num_workers=num_workers)
    valid_loader = DataLoader(valid, batch_size=32, shuffle=False,
↪num_workers=num_workers)
    test_loader = DataLoader(test, batch_size=16, shuffle=False,
↪num_workers=num_workers)

    for epoch in range(epochs):
        start_time = timer()
        correct = 0
        train_losses = 0

        # Sets the model in training mode.
        model.train()

        for images, csv_data, labels in tqdm(train_loader, desc="Training"):

```

```

        images = torch.tensor(images, device=device, dtype=torch.
↪float32)

        csv_data = torch.tensor(csv_data, device=device, dtype=torch.
↪float32)

        labels = torch.tensor(labels, device=device, dtype=torch.
↪float32)

        optimizer.zero_grad()
        out = model(images, csv_data)
        loss = criterion(out, labels.unsqueeze(1))
        loss.backward()
        optimizer.step()
        train_losses += loss.item()
        train_preds = torch.round(torch.sigmoid(out)) # 0 and 1
        correct += (train_preds.cpu() == labels.cpu().unsqueeze(1)).
↪sum().item()

        # Compute Train Accuracy
        train_acc = correct / len(train_index)

        # Set the model in evaluation mode
        model.eval()

        # Create matrix to store evaluation predictions
        valid_preds = torch.zeros(size = (len(valid_index), 1),
↪device=device, dtype=torch.float32)

        # Disables gradient tracking
        with torch.inference_mode():
            for k, (images, csv_data, labels) in
↪enumerate(tqdm(valid_loader, desc="Validation")):
                images = torch.tensor(images, device=device, dtype=torch.
↪float32)

                csv_data = torch.tensor(csv_data, device=device,
↪dtype=torch.float32)

                labels = torch.tensor(labels, device=device, dtype=torch.
↪float32)

                out = model(images, csv_data)
                pred = torch.sigmoid(out)
                valid_preds[k*images.shape[0] : k*images.shape[0] + images.
↪shape[0]] = pred

        # Compute accuracy and roc
        valid_acc = accuracy_score(valid_data['target'].values, torch.
↪round(valid_preds.cpu()))

```

```

        valid_roc = roc_auc_score(valid_data['target'].values,
↪valid_preds.cpu())

        # Compute time on Train + Eval
        duration = str(datetime.timedelta(seconds=timer() -
↪start_time))[:7]

        # PRINT INFO
        # Append to file
        with open(f"logs.txt", 'a+') as f:
            print('Time Taken:{} | Epoch: {}/{} | Loss: {:.4} | Train_
↪Acc: {:.3} | Valid Acc: {:.3} | ROC: {:.3}'.\
                format(duration, epoch+1, epochs, train_losses, train_acc,
↪valid_acc, valid_roc), file=f)
            # Print to console
            print('Time Taken:{} | Epoch: {}/{} | Loss: {:.4} | Train Acc:
↪{:.3} | Valid Acc: {:.3} | ROC: {:.3}'.\
                format(duration, epoch+1, epochs, train_losses, train_acc,
↪valid_acc, valid_roc))

        # Update scheduler and save model
        scheduler.step(valid_roc)
        # Update best_roc
        if not best_roc: # If best_roc = None
            best_roc = valid_roc
            torch.save(model.state_dict(),
↪
↪f"Fold{fold+1}_Epoch{epoch+1}_ValidAcc_{valid_acc:.3f}_ROC_{valid_roc:.3f}.
↪pth")

            continue
        if valid_roc > best_roc:
            best_roc = valid_roc
            # Reset patience
            patience_f = patience
            torch.save(model.state_dict(),
↪
↪f"Fold{fold+1}_Epoch{epoch+1}_ValidAcc_{valid_acc:.3f}_ROC_{valid_roc:.3f}.
↪pth")

        else:
            # Decrease patience
            patience_f = patience_f - 1
            if patience_f == 0:
                with open(f"logs.txt", 'a+') as f:

```

```

        print('Early stopping | Best ROC: {}'.\
              format(best_roc), file=f)
    print('Early stopping | Best ROC: {}'.\
          format(best_roc))
    break

    # Choose model with best r oc in this fold
    best_model_path = os.getcwd()+ os.sep+ [file for file in os.listdir(os.
↪getcwd()) if str(round(best_roc, 3)) in file and 'Fold'+str(fold+1) in_
↪file][0]

    model = MelanomaClassifier(output_size = output_size, no_cols=no_cols).
↪to(device)

    model.load_state_dict(torch.load(best_model_path))
    # Set the model in evaluation mode
    model.eval()

    with torch.inference_mode():
        # Predicting again on Validation data to get preds for out of fold
        valid_preds = torch.zeros(size = (len(valid_index), 1),_
↪device=device, dtype=torch.float32)

        for k, (images, csv_data, _) in_
↪enumerate(tqdm(valid_loader, desc="OOF Predictions")):
            images = torch.tensor(images, device=device, dtype=torch.
↪float32)

            csv_data = torch.tensor(csv_data, device=device, dtype=torch.
↪float32)

            out = model(images, csv_data)
            pred = torch.sigmoid(out)
            valid_preds[k*images.shape[0] : k*images.shape[0] + images.
↪shape[0]] = pred

        # Save info to OOF
        oof[valid_index] = valid_preds.cpu().numpy()

        # Initialize variables to store metrics
        total_true_positive = 0
        total_true_negative = 0
        total_false_positive = 0
        total_false_negative = 0
        total_predictions = 0

        for i in range(TTA):

```

```

        for k, (images, csv_data, labels) in
↪enumerate(tqdm(test_loader, desc="Testing")):
            images = torch.tensor(images, device=device, dtype=torch.
↪float32)
            csv_data = torch.tensor(csv_data, device=device,
↪dtype=torch.float32)
            labels = torch.tensor(labels, device=device, dtype=torch.
↪float32)

            out = model(images, csv_data)
            test_preds = torch.round(torch.sigmoid(out)) # Convert
↪logits to predictions

            preds = test_preds.cpu().numpy().flatten()
            true_labels = labels.cpu().numpy().flatten()

            # Update confusion matrix
            tn, fp, fn, tp = confusion_matrix(true_labels, preds,
↪labels=[0, 1]).ravel()
            total_true_positive += tp
            total_true_negative += tn
            total_false_positive += fp
            total_false_negative += fn

            # Update total predictions count
            total_predictions += len(preds)

total_true_positive = round(total_true_positive / TTA)
total_true_negative = round(total_true_negative / TTA)
total_false_positive = round(total_false_positive / TTA)
total_false_negative = round(total_false_negative / TTA)
total_predictions = round(total_predictions / TTA)

# Calculate metrics
accuracy = (total_true_positive + total_true_negative) /
↪total_predictions
sensitivity = total_true_positive / (total_true_positive +
↪total_false_negative)
specificity = total_true_negative / (total_true_negative +
↪total_false_positive)

results.append({
    'Fold': fold+1,
    'Accuracy': accuracy,

```

```

        'True Positive': total_true_positive,
        'True Negative': total_true_negative,
        'False Positive': total_false_positive,
        'False Negative': total_false_negative,
        'Sensitivity': sensitivity,
        'Specificity': specificity
    })
    # Clear memory
    del train, valid, train_loader, valid_loader, images, labels
    # Garbage collector
    gc.collect()

    return results

```

```

[44]: model = MelanomaClassifier(output_size = output_size, no_cols=no_cols).
      ↪to(device)
      x=train_folds(model = model)

```

Loaded pretrained weights for efficientnet-b2

----- Fold: 1 -----

Training: 100%| | 678/678 [01:08<00:00, 9.95it/s]

Validation: 100%| | 136/136 [00:05<00:00, 27.11it/s]

Time Taken:0:01:13 | Epoch: 1/30 | Loss: 400.3 | Train Acc: 0.66 | Valid Acc: 0.747 | ROC: 0.893

Training: 100%| | 678/678 [01:09<00:00, 9.80it/s]

Validation: 100%| | 136/136 [00:04<00:00, 32.42it/s]

Time Taken:0:01:13 | Epoch: 2/30 | Loss: 372.7 | Train Acc: 0.691 | Valid Acc: 0.846 | ROC: 0.875

Training: 100%| | 678/678 [01:09<00:00, 9.76it/s]

Validation: 100%| | 136/136 [00:04<00:00, 31.93it/s]

Time Taken:0:01:13 | Epoch: 3/30 | Loss: 355.8 | Train Acc: 0.711 | Valid Acc: 0.904 | ROC: 0.852

Training: 100%| | 678/678 [01:09<00:00, 9.74it/s]

Validation: 100%| | 136/136 [00:04<00:00, 32.17it/s]

Time Taken:0:01:13 | Epoch: 4/30 | Loss: 338.4 | Train Acc: 0.729 | Valid Acc: 0.882 | ROC: 0.876

Training: 100%| | 678/678 [01:09<00:00, 9.77it/s]

Validation: 100%| | 136/136 [00:04<00:00, 31.76it/s]

Time Taken:0:01:13 | Epoch: 5/30 | Loss: 328.9 | Train Acc: 0.736 | Valid Acc: 0.916 | ROC: 0.875

Early stopping | Best ROC: 0.8925800067703044

Loaded pretrained weights for efficientnet-b2

```

OOF Predictions: 100%|          | 136/136 [00:04<00:00, 32.20it/s]
Testing: 100%|          | 407/407 [00:08<00:00, 45.36it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.52it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.51it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.50it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.10it/s]

----- Fold: 2 -----

Training: 100%|          | 678/678 [01:08<00:00, 9.84it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 31.59it/s]

Time Taken:0:01:13 | Epoch: 1/30 | Loss: 370.3 | Train Acc: 0.696 | Valid Acc:
0.825 | ROC: 0.907

Training: 100%|          | 678/678 [01:09<00:00, 9.82it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.76it/s]

Time Taken:0:01:13 | Epoch: 2/30 | Loss: 355.7 | Train Acc: 0.712 | Valid Acc:
0.812 | ROC: 0.895

Training: 100%|          | 678/678 [01:09<00:00, 9.80it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 31.54it/s]

Time Taken:0:01:13 | Epoch: 3/30 | Loss: 342.1 | Train Acc: 0.728 | Valid Acc:
0.877 | ROC: 0.892

Training: 100%|          | 678/678 [01:09<00:00, 9.80it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.29it/s]

Time Taken:0:01:13 | Epoch: 4/30 | Loss: 325.6 | Train Acc: 0.741 | Valid Acc:
0.882 | ROC: 0.89

Training: 100%|          | 678/678 [01:09<00:00, 9.80it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.19it/s]

Time Taken:0:01:13 | Epoch: 5/30 | Loss: 317.6 | Train Acc: 0.744 | Valid Acc:
0.907 | ROC: 0.897
Early stopping | Best ROC: 0.9069548872180452
Loaded pretrained weights for efficientnet-b2

OOF Predictions: 100%|          | 136/136 [00:04<00:00, 32.10it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.18it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.14it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.25it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.87it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.17it/s]

----- Fold: 3 -----

Training: 100%|          | 678/678 [01:07<00:00, 10.05it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 30.68it/s]

Time Taken:0:01:11 | Epoch: 1/30 | Loss: 356.5 | Train Acc: 0.716 | Valid Acc:
0.858 | ROC: 0.915

```

```

Training: 100%|          | 678/678 [01:07<00:00, 10.04it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.42it/s]

Time Taken:0:01:11 | Epoch: 2/30 | Loss: 342.0 | Train Acc: 0.729 | Valid Acc:
0.835 | ROC: 0.913

Training: 100%|          | 678/678 [01:07<00:00, 10.07it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.10it/s]

Time Taken:0:01:11 | Epoch: 3/30 | Loss: 329.1 | Train Acc: 0.74 | Valid Acc:
0.895 | ROC: 0.896

Training: 100%|          | 678/678 [01:07<00:00, 10.03it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.00it/s]

Time Taken:0:01:11 | Epoch: 4/30 | Loss: 315.1 | Train Acc: 0.749 | Valid Acc:
0.879 | ROC: 0.893

Training: 100%|          | 678/678 [01:07<00:00, 10.05it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.16it/s]

Time Taken:0:01:11 | Epoch: 5/30 | Loss: 303.0 | Train Acc: 0.757 | Valid Acc:
0.916 | ROC: 0.888
Early stopping | Best ROC: 0.9147576830608699
Loaded pretrained weights for efficientnet-b2

OOF Predictions: 100%|          | 136/136 [00:04<00:00, 32.05it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.85it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.85it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 51.92it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.01it/s]
Testing: 100%|          | 407/407 [00:07<00:00, 52.02it/s]

----- Fold: 4 -----

Training: 100%|          | 678/678 [01:07<00:00, 10.08it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.32it/s]

Time Taken:0:01:11 | Epoch: 1/30 | Loss: 349.0 | Train Acc: 0.723 | Valid Acc:
0.885 | ROC: 0.945

Training: 100%|          | 678/678 [01:07<00:00, 10.06it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.46it/s]

Time Taken:0:01:11 | Epoch: 2/30 | Loss: 332.4 | Train Acc: 0.74 | Valid Acc:
0.875 | ROC: 0.931

Training: 100%|          | 678/678 [01:07<00:00, 10.07it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 32.11it/s]

Time Taken:0:01:11 | Epoch: 3/30 | Loss: 327.7 | Train Acc: 0.742 | Valid Acc:
0.899 | ROC: 0.924

Training: 100%|          | 678/678 [01:07<00:00, 10.04it/s]
Validation: 100%|        | 136/136 [00:04<00:00, 31.93it/s]

```



Time Taken:0:01:11 | Epoch: 4/30 | Loss: 312.7 | Train Acc: 0.751 | Valid Acc: 0.903 | ROC: 0.928

Training: 100%| | 678/678 [01:07<00:00, 10.05it/s]  
Validation: 100%| | 136/136 [00:04<00:00, 31.74it/s]

Time Taken:0:01:11 | Epoch: 5/30 | Loss: 303.6 | Train Acc: 0.763 | Valid Acc: 0.925 | ROC: 0.927

Early stopping | Best ROC: 0.9446110641055645  
Loaded pretrained weights for efficientnet-b2

OOF Predictions: 100%| | 136/136 [00:04<00:00, 31.97it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 52.36it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 52.29it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 52.24it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 52.44it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 52.50it/s]

----- Fold: 5 -----

Training: 100%| | 678/678 [01:07<00:00, 10.08it/s]  
Validation: 100%| | 136/136 [00:04<00:00, 32.10it/s]

Time Taken:0:01:11 | Epoch: 1/30 | Loss: 332.1 | Train Acc: 0.733 | Valid Acc: 0.867 | ROC: 0.947

Training: 100%| | 678/678 [01:06<00:00, 10.12it/s]  
Validation: 100%| | 136/136 [00:04<00:00, 32.60it/s]

Time Taken:0:01:11 | Epoch: 2/30 | Loss: 326.1 | Train Acc: 0.739 | Valid Acc: 0.853 | ROC: 0.938

Training: 100%| | 678/678 [01:06<00:00, 10.17it/s]  
Validation: 100%| | 136/136 [00:04<00:00, 32.63it/s]

Time Taken:0:01:10 | Epoch: 3/30 | Loss: 319.2 | Train Acc: 0.745 | Valid Acc: 0.888 | ROC: 0.935

Training: 100%| | 678/678 [01:06<00:00, 10.15it/s]  
Validation: 100%| | 136/136 [00:04<00:00, 32.75it/s]

Time Taken:0:01:10 | Epoch: 4/30 | Loss: 300.7 | Train Acc: 0.759 | Valid Acc: 0.917 | ROC: 0.944

Training: 100%| | 678/678 [01:06<00:00, 10.16it/s]  
Validation: 100%| | 136/136 [00:04<00:00, 32.93it/s]

Time Taken:0:01:10 | Epoch: 5/30 | Loss: 292.8 | Train Acc: 0.767 | Valid Acc: 0.935 | ROC: 0.946

Early stopping | Best ROC: 0.9474052432444534  
Loaded pretrained weights for efficientnet-b2

OOF Predictions: 100%| | 136/136 [00:04<00:00, 32.86it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 53.09it/s]  
Testing: 100%| | 407/407 [00:07<00:00, 52.65it/s]

```

Testing: 100%| | 407/407 [00:07<00:00, 52.95it/s]
Testing: 100%| | 407/407 [00:07<00:00, 52.73it/s]
Testing: 100%| | 407/407 [00:07<00:00, 52.87it/s]

----- Fold: 6 -----

Training: 100%| | 678/678 [01:06<00:00, 10.14it/s]
Validation: 100%| | 136/136 [00:04<00:00, 32.37it/s]

Time Taken:0:01:11 | Epoch: 1/30 | Loss: 323.9 | Train Acc: 0.748 | Valid Acc:
0.926 | ROC: 0.937

Training: 100%| | 678/678 [01:06<00:00, 10.13it/s]
Validation: 100%| | 136/136 [00:04<00:00, 32.69it/s]

Time Taken:0:01:11 | Epoch: 2/30 | Loss: 310.8 | Train Acc: 0.755 | Valid Acc:
0.922 | ROC: 0.928

Training: 100%| | 678/678 [01:06<00:00, 10.16it/s]
Validation: 100%| | 136/136 [00:04<00:00, 32.61it/s]

Time Taken:0:01:10 | Epoch: 3/30 | Loss: 305.1 | Train Acc: 0.757 | Valid Acc:
0.882 | ROC: 0.924

Training: 100%| | 678/678 [01:06<00:00, 10.14it/s]
Validation: 100%| | 136/136 [00:04<00:00, 32.41it/s]

Time Taken:0:01:11 | Epoch: 4/30 | Loss: 291.5 | Train Acc: 0.771 | Valid Acc:
0.928 | ROC: 0.92

Training: 100%| | 678/678 [01:06<00:00, 10.13it/s]
Validation: 100%| | 136/136 [00:04<00:00, 32.55it/s]

Time Taken:0:01:11 | Epoch: 5/30 | Loss: 280.8 | Train Acc: 0.778 | Valid Acc:
0.942 | ROC: 0.926
Early stopping | Best ROC: 0.9372480477663404
Loaded pretrained weights for efficientnet-b2

OOF Predictions: 100%| | 136/136 [00:04<00:00, 32.70it/s]
Testing: 100%| | 407/407 [00:07<00:00, 52.85it/s]
Testing: 100%| | 407/407 [00:07<00:00, 52.56it/s]
Testing: 100%| | 407/407 [00:07<00:00, 52.89it/s]
Testing: 100%| | 407/407 [00:07<00:00, 53.19it/s]
Testing: 100%| | 407/407 [00:07<00:00, 52.97it/s]

```

```
[45]: x = pd.DataFrame(x)
```

```
[46]: x.mean(axis=0)
```

```
[46]: Fold          3.500000
      Accuracy      0.702277
      True Positive  80.833333
      True Negative 4484.666667
```

```
False Positive    1900.333333
False Negative    35.166667
Sensitivity        0.696839
Specificity        0.702375
dtype: float64
```

```
[47]: print('ROC: {:.3f}'.format(roc_auc_score(train_df['target'], oof)))
      oof_1 = oof
```

```
ROC: 0.922
```

```
[48]: oof_1[oof_1 >= 0.5] = 1
      oof_1[oof_1 < 0.5] = 0
```

```
[49]: oof
```

```
[49]: array([[0.],
            [0.],
            [1.],
            ...,
            [0.],
            [0.],
            [0.]])
```

```
[50]: oof_1
```

```
[50]: array([[0.],
            [0.],
            [1.],
            ...,
            [0.],
            [0.],
            [0.]])
```

```
[51]: print('Out of Fold Predictions:\n',confusion_matrix(train_df['target'], oof_1))
```

```
Out of Fold Predictions:
[[21777  3794]
 [   68   391]]
```

```
[52]: print(classification_report(train_df['target'], oof_1))
```

	precision	recall	f1-score	support
0	1.00	0.85	0.92	25571
1	0.09	0.85	0.17	459
accuracy			0.85	26030

macro avg	0.55	0.85	0.54	26030
weighted avg	0.98	0.85	0.91	26030


```
[53]: # Create Confusion Matrix
cm = confusion_matrix(train_df['target'], oof_1)

# Define class labels
classes = ['Benign', 'Malignant']

# Plot confusion matrix
plt.figure(figsize=(8, 6))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion matrix')
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

# Normalize confusion matrix
cm_norm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], 'd') + '\n(' + format(cm_norm[i, j], '.2f') +
        '+ ')',
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()
```



[ ]:

[ ]: